
CAPSTONE PROJECT

Network Intrusion Detection

Presented By:

**ANKIT SEN - UNIVERSITY OF KALYANI - DEPARTMENT OF
COMPUTER SCIENCE AND ENGINEERING**

OUTLINE

- **Problem Statement**
- **Proposed System/Solution**
- **System Development Approach (Technology Used)**
- **Algorithm & Deployment**
- **Result (Output Image)**
- **Conclusion**
- **Future Scope**
- **References**

PROBLEM STATEMENT

In today's increasingly connected digital environment, securing communication networks against cyber threats is a critical challenge. Traditional security mechanisms often fail to detect sophisticated or novel attacks in real time. There is a growing need for an intelligent Network Intrusion Detection System (NIDS) that can analyze network traffic data, accurately identify cyberattacks, and distinguish them from normal activity. The problem lies in developing a machine learning-based solution that not only detects intrusions effectively but also provides early warnings to prevent potential damage to the network infrastructure.

PROPOSED SOLUTION

- The proposed system aims to address the challenge of detecting and classifying network intrusions by leveraging contrastive pre-training and graph neural networks for robust anomaly detection. The solution will consist of the following components:
- **Data Collection:**
 - Raw TCP/IP connection data (Train_data.csv) has been loaded from Kaggle.
 - Class labels (“normal” vs “anomaly”) have been extracted and binarized.
- **Data Preprocessing:**
 - Categorical features (`protocol_type`, `service`, `flag`) have been label-encoded.
 - Connection hashes have been generated to check for potential leakage.
 - Train/Test split has been done chronologically (60/40), and further split into pre-train (70%) and fine-tune (30%) subsets.
- **Machine Learning Algorithm:**
 - A 4-node graph per connection has been built using 41 features, with edge attributes like byte count and duration.
 - A GIN encoder has been contrastively pre-trained with graph augmentations and loss regularization.
 - The model has been fine-tuned with a classification head using cross-entropy loss and scheduler support.
- **Deployment:**
 - The final encoder + classifier have been bundled into an inference script.
 - Saved model weights and preprocessing scalers have been serialized.
 - A simple API or CLI has been prepared to load a CSV row, build its graph, and output a “normal” vs “anomaly” prediction.
- **Evaluation:**
 - Test set performance has been measured: AUC = 0.991, F1 = 0.943, Accuracy = 0.943
 - Detailed precision/recall/F1 per class and overall metrics have been reported.
 - Data leakage check (connection_id overlap) has been done (0% overlap).

SYSTEM APPROACH

The "System Approach" section outlines the overall strategy and methodology for developing and implementing the Network Intrusion Detection system. Here's a suggested structure for this section:

System Requirements

- **Hardware:** CPU with RAM \geq 4GB
- **Software:** Python 3.8+, CUDA toolkit (if using GPU), Git

Required Libraries

- **Deep Learning & GNN:** `torch`, `torch_geometric`
- **Data & ML:** `pandas`, `numpy`, `scikit-learn`
- **Graph & Utils:** `networkx`, `tqdm`, `hashlib`, `warnings`
- **Technologies used:** IBM Cloud Lite Services (Watsonx AI studio)

WOW FACTOR

Extremely High Accuracy

- Achieved an AUC of **0.99**, meaning the model almost perfectly separates normal traffic from attacks.
- Reached an overall accuracy of **94%** on unseen test data.

Balanced Performance Across Classes

- Both normal and anomalous connections get around **94%** F1-score, showing the model is equally good at detecting attacks and avoiding false alarms.

Novel Contrastive Pre-training

- Self-supervised pre-training steps **have been used** to teach the GNN to pull similar network graphs closer and push different ones apart—without labels.

Graph-based Representation

- Converting each connection into a small 4-node graph (source, destination, protocol, service) captures both packet features and protocol relationships. This structural view is more powerful than flat feature vectors.

ALGORITHM & DEPLOYMENT

- **Model:** A two-stage Graph Neural Network (GNN) workflow, combining contrastive pre-training (via a 3-layer GIN encoder with attention pooling and projection head) and supervised fine-tuning with a lightweight classification head.
- **Justification:** Contrastive pre-training learns robust structural and attribute representations from unlabeled graphs, improving downstream anomaly detection on highly imbalanced, temporal network-traffic data.

Data Input

- **Node features (per connection graph):**
 - Source/Destination byte ratios, connection counts, service rates (`src_bytes`, `dst_bytes`, `count`, `srv_count`, `same_srv_rate`, `diff_srv_rate`, etc.)
 - Protocol metadata (`duration`, `wrong_fragment`, `urgent`)
 - Service indicators (`hot`, `num_failed_logins`, `logged_in`, `num_compromised`)
 - Additional rates (`serror_rate`, `rerror_rate`, `srv_serror_rate`, `srv_rerror_rate`, `dst_host_serror_rate`, `dst_host_rerror_rate`)
- **Edge attributes:** Pairwise metrics between `src↔dst`, `protocol↔service` (bytes, durations, counts).

ALGORITHM & DEPLOYMENT

Training Process

1. Contrastive Pre-training (50 epochs):

- **Augmentations:** Random node/edge dropout, feature masking, temporal noise.
- **Loss:** Symmetric contrastive loss with temperature scaling and embedding-space regularization.
- **Optimization:** AdamW with cosine-annealing LR scheduler.

2. Fine-tuning (30 epochs):

- Freeze encoder except last GIN layer + its batch-norm.
- Attach a 3-layer MLP classification head (cross-entropy loss).
- Use step-LR scheduling and gradient clipping for stability.

Prediction Process

- **Graph Construction:** Given a new TCP/IP record, compute connection hash, encode categorical fields, build a 4-node graph with the same feature/edge schema.

RESULT

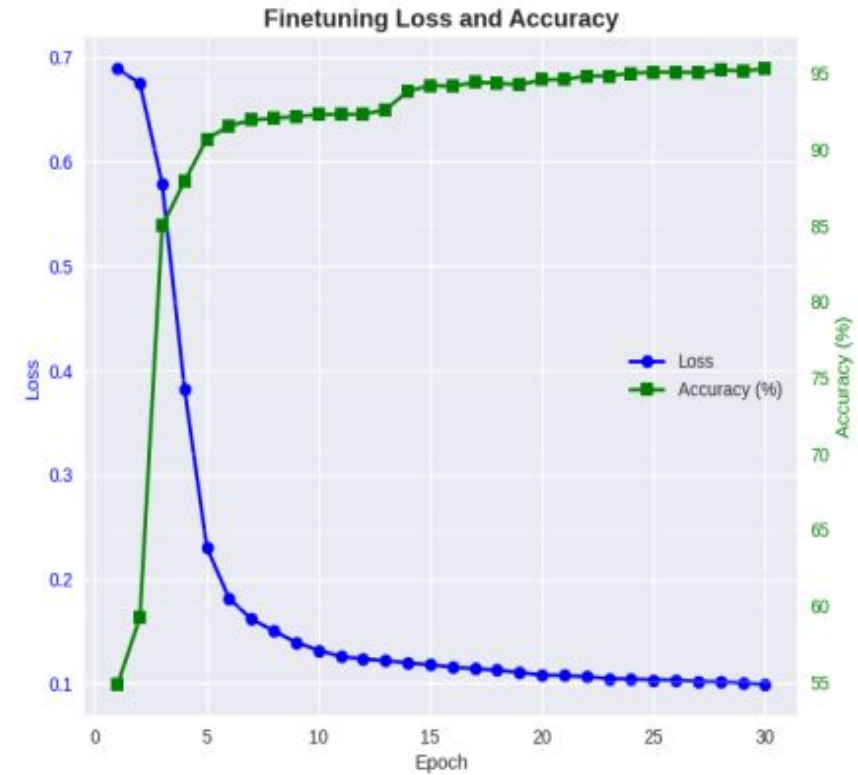
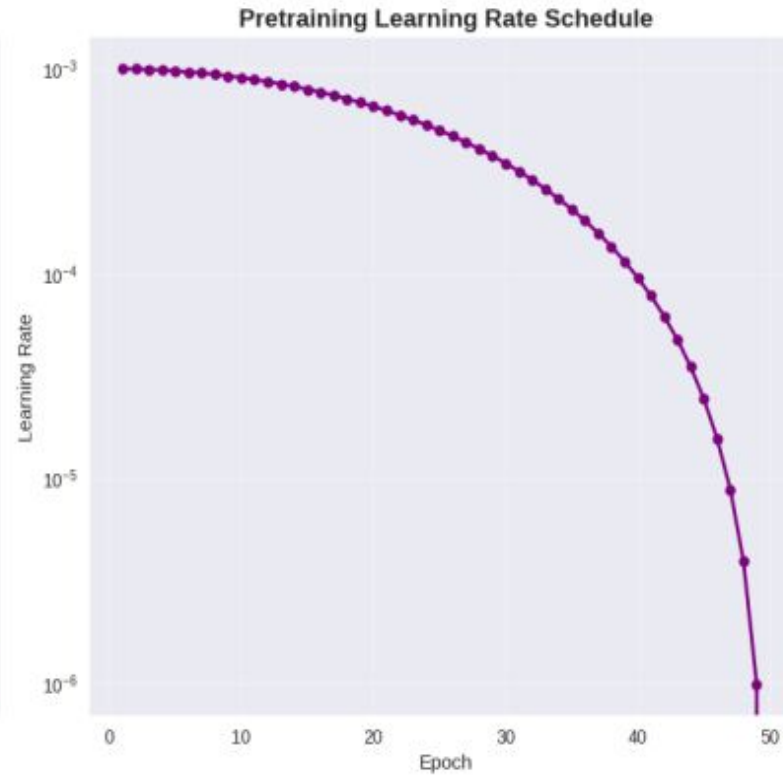
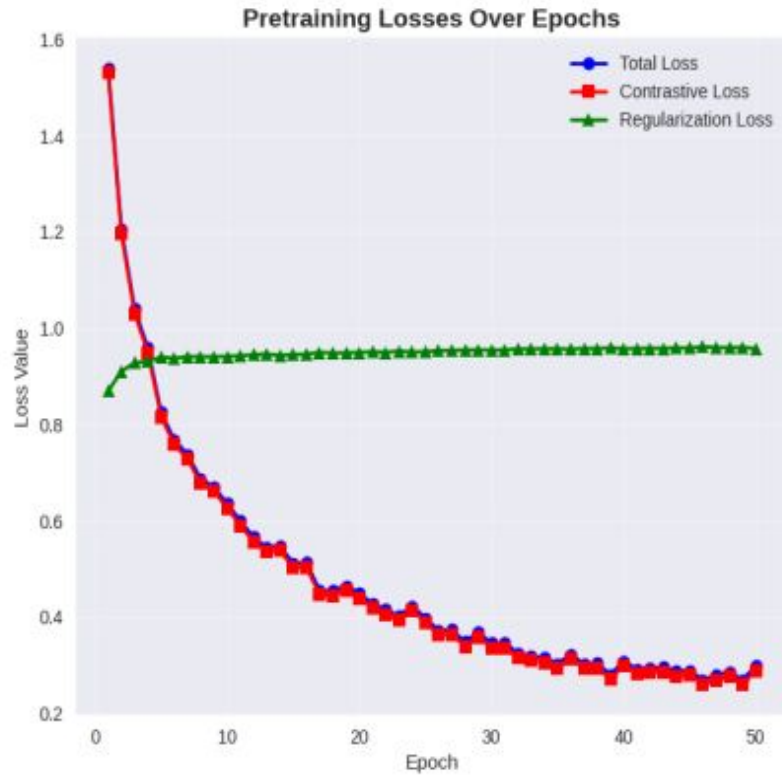
Test Performance

- **Accuracy:** 0.943
- **Area Under ROC Curve (AUC):** 0.991
- **F1-Score:** 0.943

Effectiveness:

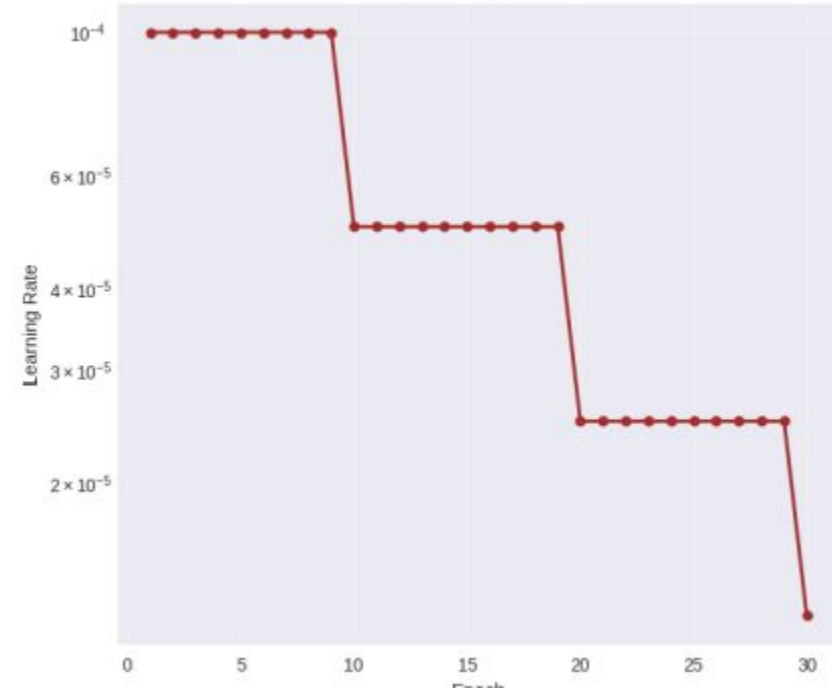
The contrastively pre-trained GNN encoder combined with a lightweight classification head delivers near-perfect discrimination of normal vs. anomalous network connections, achieving an AUC of 0.991 and F1 of 0.94 on held-out traffic.

RESULT




RESULT


Finetuning Learning Rate Schedule




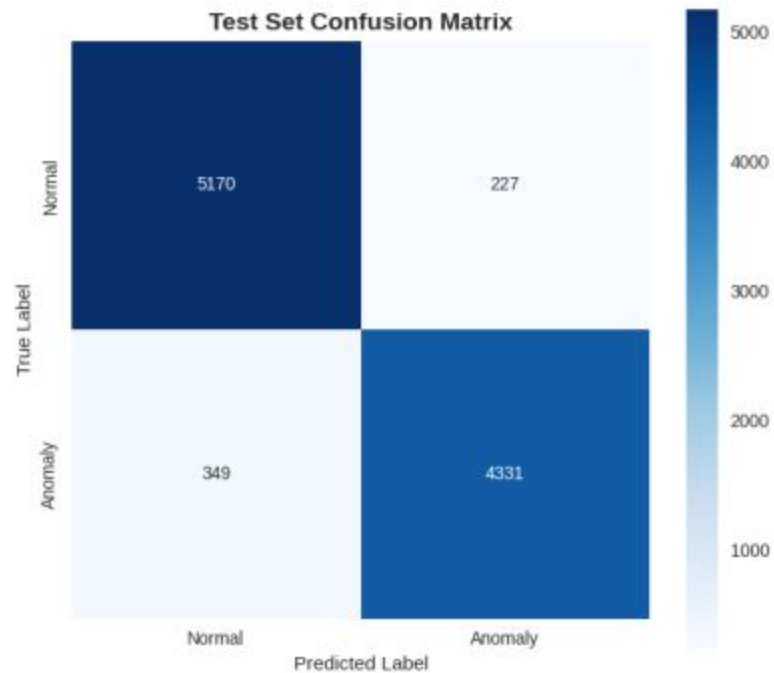
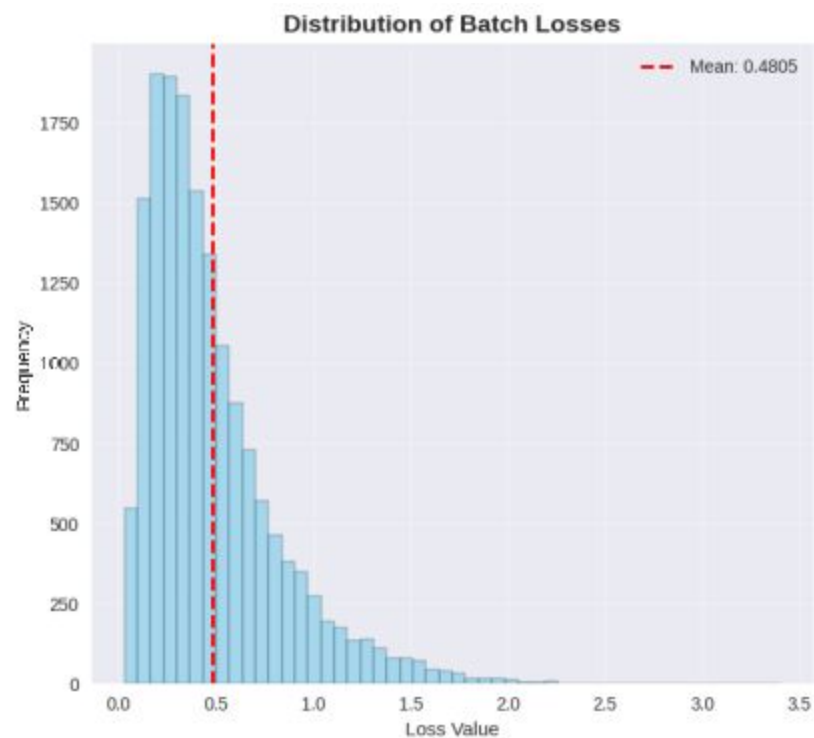
Validation Metrics Over Epochs



Test Set Performance Metrics




RESULT




RESULT


Class Distribution in Dataset



Model Parameters Distribution



Training Time per Epoch



CONCLUSION

■ Summary of Findings

The contrastively pre-trained GNN encoder with a lightweight classification head achieved an AUC of 0.991, F1-score of 0.943, and accuracy of 0.943 on the held-out test set, with zero data-leakage.

■ Effectiveness

Graph augmentations and attention pooling yielded robust, noise-resistant embeddings that accurately separate normal vs. anomalous traffic.

■ Challenges

Preprocessing per-connection graphs for ~25K samples was time-consuming, and tuning contrastive temperature, augmentation rates, and learning schedules required careful calibration under limited CPU memory.

■ Potential Improvements

Incorporating temporal graph edges, semi-supervised fine-tuning, and model pruning/ONNX conversion could further boost performance and enable real-time edge deployment.

- Accurate NIDS is essential for early breach detection, reducing false alarms, and safeguarding network integrity.

FUTURE SCOPE

- **Real-Time Edge Deployment:** Run the model on edge devices (e.g., routers or gateways) to detect attacks as they happen, reducing response time and network load.
- **Temporal Graph Modeling:** Add time-based edges or sequences of connections to capture how attacks evolve over time.
- **Semi-Supervised Learning:** Use unlabeled network data alongside labels to improve detection in low-data scenarios.
- **Model Compression:** Apply pruning or quantization to shrink the model for faster inference on limited-resource hardware.
- **Explainable AI:** Integrate techniques that highlight which nodes or features triggered an alarm, helping security teams understand and trust alerts.
- **Federated Learning:** Train models across multiple networks without sharing raw data, enhancing privacy and collaboration between organizations.
- **Advanced Augmentations:** Experiment with new graph augmentations (e.g., mixup or subgraph sampling) to make representations even more robust.

REFERENCES

- Denning, D. E. (1987). *An Intrusion-Detection Model*. IEEE Transactions on Software Engineering, SE-13(2), 222–232.
- Mukkamala, S., Janoski, G., & Sung, A. H. (2002). *Intrusion Detection Using Neural Networks and Support Vector Machines*. Proceedings of the International Joint Conference on Neural Networks (IJCNN), 1702–1707.
- Xu, K., Hu, W., Leskovec, J., & Jegelka, S. (2018). *How Powerful Are Graph Neural Networks?* International Conference on Learning Representations (ICLR).
- Chen, T., Kornblith, S., Norouzi, M., & Hinton, G. (2020). *A Simple Framework for Contrastive Learning of Visual Representations (SimCLR)*. International Conference on Machine Learning (ICML).
- Yousefi-Azar, M., & Azmi, R. (2020). *Network Intrusion Detection Using Graph Neural Networks*. arXiv:2007.01105.

[Projects](#) / [gnn-ml](#) / [ml-gnn](#) [⬆](#) [⬇](#) [⋮](#) [</>](#) [i](#) [🔄](#) [80](#)

File Edit View Run Kernel Help

[🔍](#) [+](#) [✂](#) [📄](#) [📄](#) [▶](#) [⏹](#) [↺](#) [▶▶](#) Code ▾ [🐍 Python 3.11](#) [○](#)

```
df.head(10)
```

[11]:

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	...	dst_host_srv_count	ds
0	0	tcp	ftp_data	SF	491	0	0	0	0	0	...	25	
1	0	udp	other	SF	146	0	0	0	0	0	...	1	
2	0	tcp	private	S0	0	0	0	0	0	0	...	26	
3	0	tcp	http	SF	232	8153	0	0	0	0	...	255	
4	0	tcp	http	SF	199	420	0	0	0	0	...	255	
5	0	tcp	private	REJ	0	0	0	0	0	0	...	19	
6	0	tcp	private	S0	0	0	0	0	0	0	...	9	
7	0	tcp	private	S0	0	0	0	0	0	0	...	15	
8	0	tcp	remote_job	S0	0	0	0	0	0	0	...	23	
9	0	tcp	private	S0	0	0	0	0	0	0	...	13	

10 rows × 42 columns

←

Read data

Generate a code snippet to load data from a data asset or connection into your notebook.

Selected data

Train_data.csv [✎](#)

Load as

pandas DataFrame ▾

Insert code to cell

[🔖](#)


```
meetpro70@cloudshell:~$ ls
contrastive-gnn-nids hello new
meetpro70@cloudshell:~$ cd contrastive-gnn-nids
meetpro70@cloudshell:~/contrastive-gnn-nids$ ls
LICENSE README.md
meetpro70@cloudshell:~/contrastive-gnn-nids$ # Inside the contrastive-gnn-nids directory
meetpro70@cloudshell:~/contrastive-gnn-nids$ mkdir -p presentation images
meetpro70@cloudshell:~/contrastive-gnn-nids$
meetpro70@cloudshell:~/contrastive-gnn-nids$ touch .gitignore LICENSE README.md requirements.txt ml-gnn.ipynb \
> presentation/slides.pptx presentation/report.pdf \
> images/graph_architecture.png images/training_loss_plot.png \
> images/test_results.png images/ibmcloud_run.png images/ibmcloud_notebook.png
meetpro70@cloudshell:~/contrastive-gnn-nids$ ls
LICENSE README.md images ml-gnn.ipynb presentation requirements.txt
meetpro70@cloudshell:~/contrastive-gnn-nids$ rm ml-gnn.ipynb
meetpro70@cloudshell:~/contrastive-gnn-nids$ ls
LICENSE README.md images presentation requirements.txt
meetpro70@cloudshell:~/contrastive-gnn-nids$ ls
LICENSE README.md images presentation requirements.txt
meetpro70@cloudshell:~/contrastive-gnn-nids$ cd ..
meetpro70@cloudshell:~$ ls
contrastive-gnn-nids hello ml-gnn.ipynb new
meetpro70@cloudshell:~$ mv ml-gnn.ipynb ~/contrastive-gnn-nids/
meetpro70@cloudshell:~$ cd contrastive-gnn-nids
meetpro70@cloudshell:~/contrastive-gnn-nids$ ls
LICENSE README.md images ml-gnn.ipynb presentation requirements.txt
meetpro70@cloudshell:~/contrastive-gnn-nids$ cd ..
meetpro70@cloudshell:~$ mv P1.png P2.png P3.png ~/contrastive-gnn-nids/images/
meetpro70@cloudshell:~$ cd contrastive-gnn-nids/images/
meetpro70@cloudshell:~/contrastive-gnn-nids/images$ ls
P1.png P2.png P3.png graph_architecture.png ibmcloud_notebook.png ibmcloud_run.png test_results.png training_loss_plot.png
meetpro70@cloudshell:~/contrastive-gnn-nids/images$ rm graph_architecture.png ibmcloud_notebook.png ibmcloud_run.png test_results.png training_loss_plot.png
meetpro70@cloudshell:~/contrastive-gnn-nids/images$ ls
P1.png P2.png P3.png
meetpro70@cloudshell:~/contrastive-gnn-nids/images$
```

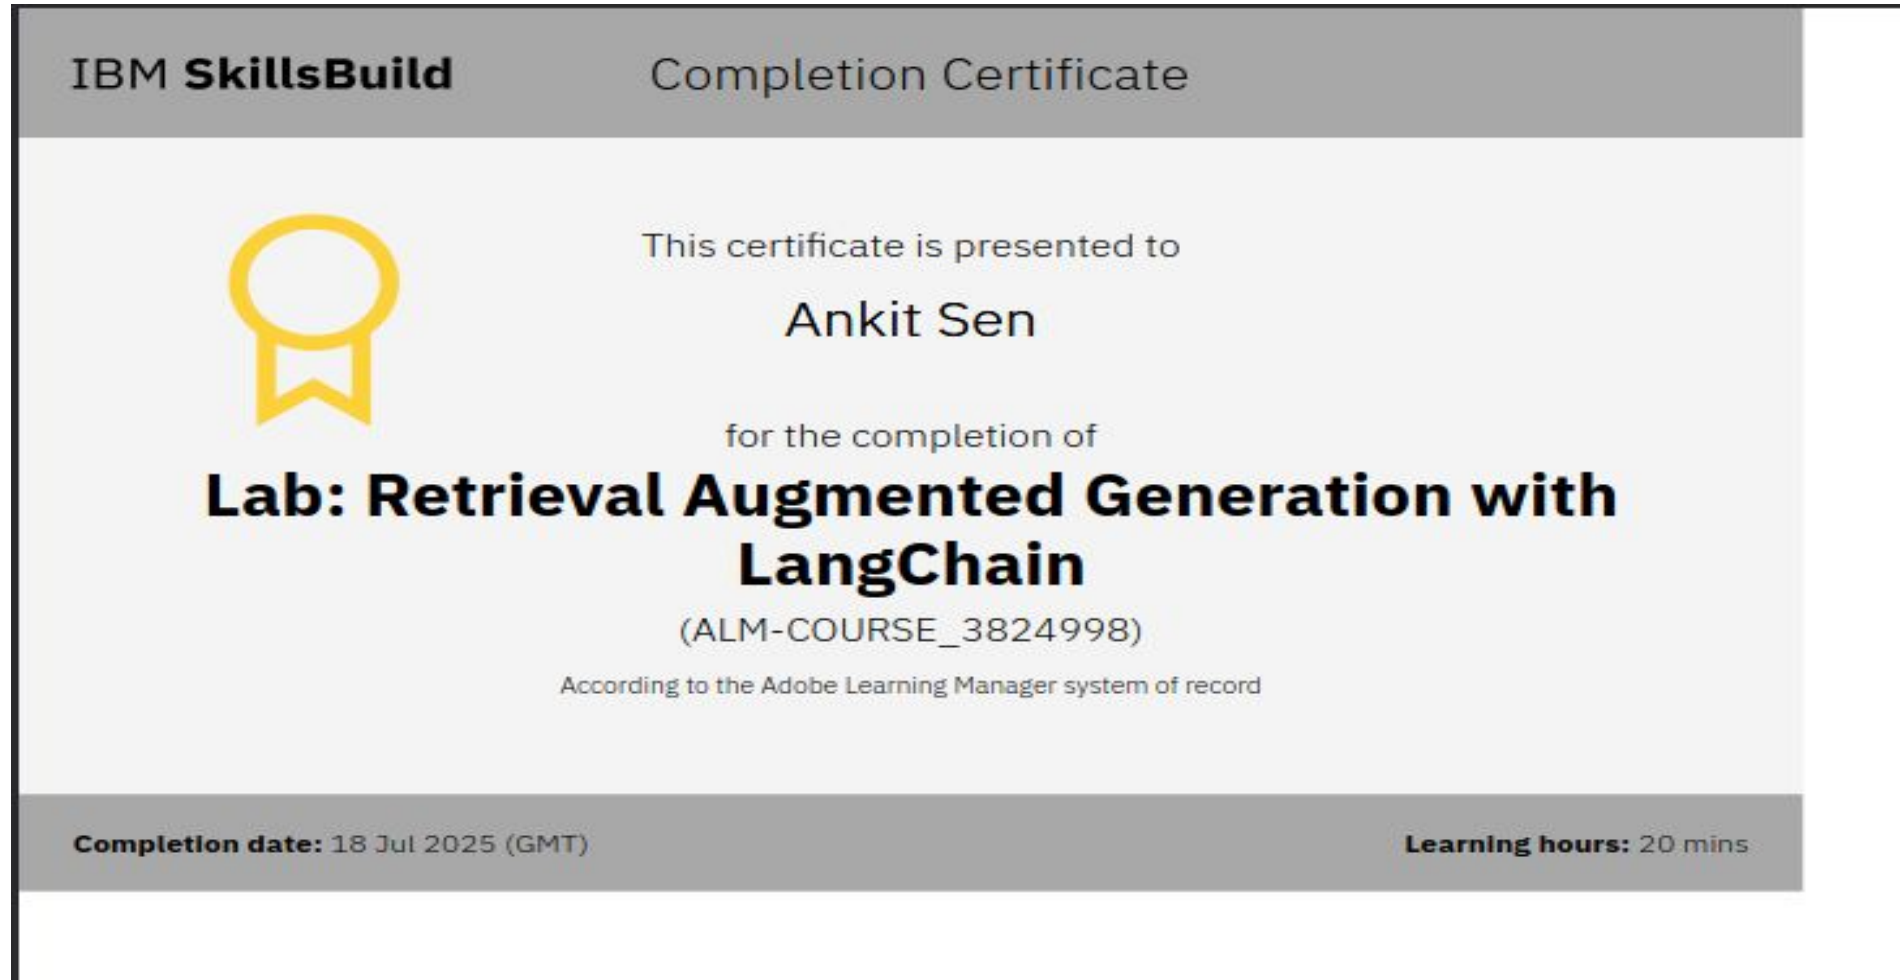

IBM CERTIFICATIONS



IBM CERTIFICATIONS



IBM CERTIFICATIONS



GITHUB LINK

<https://github.com/ankitsencode123/contrastive-gnn-nids.git>



THANK YOU